

Multi-Agent LLM Framework for 5G Network Slice Resource Management

Zheyun Zhao

Mid-Term Review

BUPT No: 2022213670 QMUL No: 221170559

March 13, 2026

Outline

Background & Motivation

Methodology

Experimental Results

Problems, Solutions & Future Work

Research Background · Existing Challenges · Motivation

5G Network Slicing creates multiple logical networks over shared infrastructure. Three slice types share limited spectrum with **competing** QoS (Quality of Service) demands:


eMBB (high throughput) / **URLLC** (ultra-low latency) / **mMTC** (massive access). Traditional rule-based and DRL (Deep Reinforcement Learning) methods struggle to unify 3GPP SLA (Service Level Agreement) constraints, operator policies, and dynamic user behaviour into an adaptive framework.

eMBB (Enhanced Mobile Broadband): high-throughput services, e.g. video streaming.

URLLC (Ultra-Reliable Low-Latency Comm.): mission-critical apps, sub-ms latency.

mMTC (Massive Machine-Type Comm.): massive IoT device connectivity.

Framework	Architecture	Multi-Agent	Negotiation	E2E Slicing
WirelessAgent [1]	Percept-Memory-Plan-Act	✗	✗	✗
CommLLM [2]	Three-component MAS	✓	✗	✗ Semantic
Qu et al. [3]	Dual-loop edge-terminal	✓	Centralised	Partial
MetaGPT [4] / ChatDev [5]	Role-based cooperative	✓	✓ Cooperative	✗ General
This Work	Competitive MAS	✓	✓ Competitive	✓



Fixed B_{total} , zero-sum inter-slice competition

Description

Gap 1 End-to-end multi-agent LLM (Large Language Model) framework for slicing remains underexplored

Gap 2 Competitive negotiation for constrained resource allocation is unaddressed by cooperative MAS (Multi-Agent System)


RQ1 Can multi-agent LLM outperform rule-based and single-agent baselines?

RQ2 Does negotiation provide structural advantages over direct allocation?

Methodology: Architecture · Two-Tier Scheduling · Negotiation

System Architecture — Agent Layer / Collaboration Layer (LangGraph) / Environment Layer


(Gymnasium):



Negotiation Flow (LangGraph StateGraph, max 3 rounds):



Two-Tier Scheduling — Bridging LLM inference latency (seconds) vs network real-time requirements (milliseconds):



ID	Method	Category	Description
M1	Fixed-Ratio	Rule-based	Fixed 5:3:2 allocation
M2	Threshold	Rule-based	SLA-triggered dynamic adjustment
M4	Single-Agent LLM	Single LLM	GPT-4 + CoT (Chain-of-Thought)
M5-PA	Multi-Agent PA	Multi LLM	Priority arbitration negotiation
M5-PC	Multi-Agent PC	Multi LLM	Proportional compromise negotiation


Experiment 1: Core Performance Comparison (RQ1)

Steady-State Scenario (100 users, 100 decision cycles $\times$ 2 seeds):

Method	SLA_{eMBB}	SLA_{URLLC}	SLA_{mMTC}	$\bar{S}$	Fairness	Lat. (ms)	Tokens
M1-Fixed	0.35 ± 0.05	0.52 ± 0.03	0.58 ± 0.07	0.48	0.893	0	0
M2-Threshold	0.31 ± 0.04	0.44 ± 0.02	0.51 ± 0.04	0.42	0.912	0	0
M4-SingleLLM	0.33 ± 0.01	0.90 ± 0.01	0.22 ± 0.05	0.48	0.720	6326	26k
M5-PA	0.21 ± 0.03	0.85 ± 0.06	0.15 ± 0.04	0.40	0.616	8065	125k
M5-PC	0.38 ± 0.03	0.76 ± 0.02	0.73 ± 0.10	0.62	0.952	8033	125k

Burst Scenario (eMBB users doubled at cycle 50):

Method	SLA_{eMBB}	SLA_{URLLC}	SLA_{mMTC}	$\bar{S}$	Fairness	Lat. (ms)	Tokens
M1-Fixed	0.21 ± 0.06	0.63 ± 0.02	0.54 ± 0.08	0.46	0.831	0	0
M2-Threshold	0.19 ± 0.03	0.38 ± 0.01	0.45 ± 0.05	0.34	0.876	0	0
M4-SingleLLM	0.24 ± 0.02	0.95 ± 0.00	0.24 ± 0.14	0.48	0.660	6063	26k
M5-PA	0.12 ± 0.11	0.89 ± 0.08	0.15 ± 0.03	0.38	0.541	7957	125k
M5-PC	0.28 ± 0.02	0.80 ± 0.06	0.68 ± 0.06	0.59	0.924	7887	125k






Key Findings (RQ1):

- M5-PC outperforms all baselines in overall SLA ($\bar{S}=0.62$) with Fairness=0.952, far above M4 (0.720) and M5-PA (0.616)
- M5-PA over-prioritises URLLC at the expense of eMBB/mMTC
- M4 tends toward extreme allocation: very high URLLC (0.90) but very low mMTC (0.22)
- Rule methods have zero latency but cannot reason about SLA semantics
- **Proportional compromise is the best strategy for balancing SLA and fairness**

Experiment 2: Negotiation Mechanism Validation (RQ2)

Method	Scenario	SLA _{eMBB}	SLA _{URLLC}	SLA _{mMTC}	$\bar{S}$	Fairness
M5-PA	steady	0.21	0.85	0.15	0.40	0.616
M5-PA	burst	0.12	0.89	0.15	0.38	0.541
M5-PC	steady	0.38	0.76	0.73	0.62	0.952
M5-PC	burst	0.28	0.80	0.68	0.59	0.924
M5-NoNeg	steady	0.56	0.49	0.14	0.39	0.810
M5-NoNeg	burst	0.32	0.71	0.36	0.46	0.839



Key Findings (RQ2) — Structural Advantage of Negotiation:

- **M5-PC vs M5-NoNeg (steady):** $\bar{S} + 59\%$ ($0.39 \rightarrow 0.62$), Fairness $+18\%$ ($0.81 \rightarrow 0.95$)
- **M5-PC vs M5-PA:** Proportional compromise outperforms priority arbitration across all metrics, especially mMTC SLA (0.73 vs 0.15)
- **M5-NoNeg** uses LLM but lacks iterative negotiation, resulting in imbalanced allocation (eMBB=0.56, mMTC=0.14)
- **M5-PA** hard priority causes URLLC to monopolise resources ($0.85+$); low-priority slices are systematically under-served
- **Conclusion:** The negotiation protocol itself (not merely more LLM calls) drives significant SLA and fairness improvements

Problems & Solutions · Future Work

Identified Problems

- **MAS latency incompatible with real-time control.** Multi-agent negotiation takes ~ 8 s per step, far exceeding the millisecond-level response time required by network control loops.
- **Statistical significance.** Current experiments use only 2 random seeds with 100 decision cycles per episode, which limits the robustness of conclusions.

Proposed Solutions

- **Strategic MAS + tactical rule executor** (*in progress*): Reposition MAS as a **strategic layer** that predicts traffic trends and generates allocation *policies*, rather than making per-TTI decisions. The lower-tier rule executor operates at 10 ms granularity guided by MAS-provided policies, decoupling LLM inference from the real-time control loop.
- **Expanded experimental scale:** Increase to 5 random seeds $\times$ 500 decision cycles per episode to strengthen statistical validity.

Future Work

- **Strategic MAS + tactical rule executor:** Validate the two-tier paradigm where MAS provides medium-term allocation policies and the rule layer handles real-time execution
- **DRL baseline (M3-PPO):** Train PPO (Proximal Policy Optimisation) with Stable-Baselines3 as an additional baseline
- **RAG (Retrieval-Augmented Generation) integration:** Inject 3GPP specifications and telecom domain knowledge into agent prompts
- **Agent memory:** Leverage historical allocation patterns to improve decision quality over time

Setup: 100 MHz (NR n78), 100 users, GPT-4o-mini (temp=0), negotiation ≤ 3 rounds, Gymnasium + LangGraph.

References


- [1] Jingwen Tong et al. “WirelessAgent: Large Language Model Agents for Intelligent Wireless Networks”. In: *arXiv preprint arXiv:2505.01074* (2025).
- [2] Feibo Jiang et al. “Large Language Model Enhanced Multi-Agent Systems for 6G Communications”. In: *IEEE Wireless Communications* 31.6 (2024), pp. 48–55.
- [3] Zhiyuan Qu et al. “LLM Enabled Multi-Agent System for 6G Networks: Framework and Method of Dual-Loop Edge-Terminal Collaboration”. In: *arXiv preprint arXiv:2509.04993* (2025).
- [4] Sirui Hong et al. “MetaGPT: Meta Programming for Multi-Agent Collaborative Framework”. In: *arXiv preprint arXiv:2308.00352* (2023).
- [5] Chen Qian et al. “ChatDev: Communicative Agents for Software Development”. In: *arXiv preprint arXiv:2307.07924* (2023).

Thank You!


Appendix A: Simulation Environment Validation




Poisson arrival validation (χ^2 test)



CQI AR(1) mean-reverting evolution




Burst event impact validation



Inter-slice resource competition tradeoff


Appendix B: Experiment 1 — Comprehensive Performance Overview

Experiment 1: Comprehensive Performance Comparison



Throughput, bandwidth utilisation, fairness, decision latency, and API cost across all five methods


Appendix C: Experiment 1 — Burst Scenario SLA Trend



SLA satisfaction rate over decision cycles under burst scenario (eMBB users doubled at cycle 50)

Appendix D: Experiment 2 — Comprehensive Performance Overview

Experiment 2: Negotiation Mechanism Performance



M5-PA / M5-PC / M5-NoNeg throughput, utilisation, latency, and API cost comparison